

# Mitigating Context Pollution in Neurosymbolic Widening via Slicing and Caching

VISHAL VUNNAM, University of Colorado Boulder, USA

Static Program Analysis is a cornerstone technique in software engineering for ensuring code reliability and security. Traditional static analysis, using abstract interpretation, represents the program as a set of mathematical constraints over abstract domains. These analyzers are sound, but often imprecise due to the undecidability of program termination, leading to excessive false positives. Recent work has explored using Language Models to improve the precision of widening-based static analyzers by predicting loop invariants and abstraction heuristics. However, these approaches explode in both context size and computation time as the number of program variables increases, making them infeasible for real-world programs. In this work, we propose SLICE-ABSINT, a novel approach to static analysis that leverages Language Models to guide the analysis process while maintaining scalability. By slicing the program based on the dependencies of diverging abstract states, we dynamically prune the context window to strictly relevant code paths before querying the neural oracle. This semantic filtering prevents “context pollution,” ensuring the model focuses solely on variables affecting the widening decision.

## 1 INTRODUCTION

Software systems are becoming increasingly complex. As mission-critical backend applications migrate toward dynamic languages such as Python and JavaScript, the need for robust static analysis, for compiler optimizations, vulnerability detection, and formal verification—has never been greater. However, the dynamic features of these languages pose a fundamental challenge. Traditional static analysis techniques, such as Abstract Interpretation [3], provide soundness, the formal guarantee that the analysis will not miss any potential errors. Yet, to maintain this guarantee in the face of dynamic ambiguity, traditional analyzers must over-approximate the program’s behavior. This loss of precision results in excessive false positives, often rendering the tools impractical for developers.

Recent work, specifically AbsInt-AI [4], has demonstrated that Language Models (LMs) can play a pivotal role in supporting sound static analysis. By leveraging the contextual and semantic understanding of LMs, these systems can predict precise heap abstractions and loop invariants that traditional heuristics miss. Crucially, ABSINT-AI was the first framework to integrate LMs without sacrificing soundness, using the model solely as a heuristic oracle while relying on the underlying abstract interpreter to mathematically verify the suggestions.

However, this approach faces a critical scalability barrier. Current methods naively present the LM with the entire program state to make a local decision. As the number of program variables increases, the context window becomes flooded with irrelevant code and data. We term this phenomenon “Context Pollution.”

Context pollution leads to two failures: (1) Inefficiency, as token costs and latency explode linearly with program size rather than complexity, and (2) Imprecision, as the LM struggles to separate semantically relevant variables from noise, leading to degraded invariant predictions.

To address these challenges, we propose SLICE-ABSINT, a novel framework that integrates semantic program slicing directly into the LM-guided abstract interpretation loop.

Our approach differs fundamentally from the status quo. Instead of querying the LM with the full program state, SLICE-ABSINT analyzes the control and data dependencies of the diverging abstract state. It constructs a precise backward slice containing only the code segments that mathematically influence the variables of interest. This targeted approach mitigates context pollution by presenting the LM with a minimal, semantically potent context.

---

Author’s address: Vishal Vunnam, vivu9928@colorado.edu, University of Colorado Boulder, USA.

Furthermore, to improve runtime and reduce token consumption, we propose a caching mechanism of the graph dependa

In summary, this paper makes the following contributions:

- **Methodology:** We propose SLICE-ABSINT, a framework that leverages semantic slicing to optimize the prompt engineering of neurosymbolic widening operators.
- **Soundness:** We formalize the slicing mechanism and prove that invariants derived from a valid program slice are observational equivalents of the full program, preserving the soundness of the verification.
- **Efficiency:** We demonstrate that our targeted context reduction significantly lowers token consumption and improves invariant prediction accuracy compared to full-context baselines.

## 2 OVERVIEW

### 3 MOTIVATING EXAMPLE

### 4 METHODOLOGY: ABSTRACT INTERPRETATION AND ABSINT-AI BACKGROUND

#### 4.1 Abstract Interpretation Framework

Abstract Interpretation provides a general framework for approximating the semantics of a program. In this work, we adopt the standard conventions where concrete values  $v \in \mathcal{V}$  are mapped to abstract values  $\hat{v} \in \mathcal{A}$  via an abstraction function  $\alpha$ , and abstract values are mapped back to sets of concrete values via a concretization function  $\gamma$ .

We define our specific abstract interpreter, ABSINT, as a state transition system operating on the abstract state tuple  $\Sigma^\# = \langle H_L, H_G, \sigma \rangle$ .

- **Local Heap ( $H_L$ ):** A flow-sensitive mapping  $H_L : \text{Addr}_L \rightarrow \text{Obj}^\#$ , tracking objects allocated within the current lexical scope.
- **Global Heap ( $H_G$ ):** A flow-insensitive mapping  $H_G : \text{Addr}_G \rightarrow \text{Obj}^\#$ , capturing the “soup” of globally visible objects and shared state, crucial for modeling the asynchronous event loops typical of dynamic languages.
- **Stack ( $\sigma$ ):** A mapping  $\sigma : \text{Var} \rightarrow \text{Val}^\#$ , storing local variable bindings.

The abstract values  $\text{Val}^\#$  are defined over a lattice that includes disjoint domains for primitives and references:

$$\text{Val}^\# ::= \perp \mid \top \mid \text{Int}^\# \mid \text{Bool}^\# \mid \text{Ref}(\text{Addr})$$

where  $\text{Int}^\#$  is the Interval domain  $[l, u]$  with bounds in  $\mathbb{Z} \cup \{-\infty, +\infty\}$ .

#### 4.2 Execution Phases of ABSINT-AI

The analysis proceeds in two distinct phases to balance precision and convergence.

*Phase 1: Small-Step Execution.* For standard linear statements (assignments, arithmetic), the interpreter executes in a flow-sensitive manner.

*Phase 2: LM-Guided Summarization.* When the analyzer encounters unbounded structures, such as ‘while’ loops or recursion, standard fixed-point iteration may fail to converge within a reasonable time. Traditional static analyzers employ a widening operator ( $\nabla$ ) to force convergence, often at the cost of precision (e.g., widening  $[0, 1]$  directly to  $[0, \infty]$ ).

ABSINT-AI replaces this blind widening with a *Semantic Summarization* step. At the loop head, the analyzer invokes an Oracle  $O_{LM}$  (a Language Model, we are running benchmarks on Ollama’s Ilama3). The oracle identifies diverging variables and suggests an abstraction strategy, either merging variables into primitive abstract domains or collapsing complex objects into summary nodes.

99 However, querying  $O_{LM}$  with the full abstract state  $\Sigma^\#$  introduces *Context Pollution*. The inclusion  
 100 of irrelevant variables noise-gates the model, reducing the accuracy of the invariant prediction and  
 101 linearly increasing the cost of analysis.

## 102 103 [CONTRIBUTION 1] SEMANTIC SLICING (SAINT)

104 To mitigate context pollution, we introduce the Semantic Analysis via INcremental Truncation  
 105 (SAINT) algorithm. The core of SAINT is a backward slicing operator that filters the program  
 106 context before it reaches the LM.

### 107 108 5.1 Formal Definition of Slicing

109 I am defining the slice operator  $\mathcal{S}$  as a backwards fixed-point calculation on the Program Dependence  
 110 Graph (PDG) [5]. Given a program  $P$ , a location  $\ell$ , and a set of diverging variables  $V_{div}$  identified by  
 111 the interpreter, the slicer computes a program  $P_{slice}$  that is a dependency closure of the defined  
 112 variables.

113 The algorithm tracks a set of relevant statements  $P_{slice}$  and a set of tracked variables  $V_{trace}$   
 114 (initialized to  $V_{div}$ ). Below are inference rules defining the iteration of the program towards a fixed  
 115 point. In the appendix are more specific inference rules for various statement types.

116 *Rule 1: Data Dependency*. If a statement  $s$  defines a variable currently in the trace set, it must be  
 117 included.

$$\frac{s \notin P_{slice} \wedge (\text{Def}(s) \cap V_{trace} \neq \emptyset)}{(P_{slice}, V_{trace}) \longrightarrow (P_{slice} \cup \{s\}, V_{trace} \cup \text{Use}(s))}$$

121 This captures the flow of values. For example, if  $x \in V_{trace}$  and the statement is  $x := y + 1$ , then the  
 122 statement is added and  $y$  is added to  $V_{trace}$ .

123 *Rule 2: Control Dependency*. If a statement  $s'$  is already in the slice, and its execution is in the  
 124 scope of a control statement. That control statement must be added to the program slice, along  
 125 with its conditional variable.

$$\frac{s \notin P_{slice} \wedge (\exists s' \in P_{slice} : s' \in \text{Scope}(s))}{(P_{slice}, V_{trace}) \longrightarrow (P_{slice} \cup \{s\}, V_{trace} \cup \text{Use}(s))}$$

129 For example, if  $s'$  is inside an if (b) { ... } block and  $s'$  is in the slice, then the condition  $b$   
 130 must be added to the slice and its variables added to  $V_{trace}$ .

131 These rules can be applied iteratively, into two representation of the sub-program:

- 132 • **Dependency Graph Representation** A PDG representing the control and data dependen-  
 133 cies of the program.
- 134 • **A Rebuilt Syntactic Sub-Program** A syntactic representation of the program slice, rebuilt  
 135 from the PDG by performing a topological sort on the dependency graph.

137 Optimally, the LM might perform better with the syntactic representation. Work can and should  
 138 be done to compare the two representations.

### 139 140 5.2 The Summarization Process

141 The slicing operator is integrated into the summarization function as follows:

- 142 (1) **Identification:** The abstract interpreter detects a loop at location  $\ell$  and identifies variables  
 143  $V_{div}$  that have changed since the last iteration. Thinks that it should call summary to get  
 144 invariant for these variables.
- 145 (2) **Pruning:** Compute  $P_{slice} = \mathcal{S}(P, \ell, V_{div})$ . Or in other words, get the slice of the program  
 146 relevant to the diverging variables at the loop head.

148 (3) **Projection:** Project the current abstract state  $\Sigma^\#$  to restrict it to only the variables in  $V_{trace}$ .

149 Let  $\Sigma_{proj}^\# = \Sigma^\# \upharpoonright V_{trace}$  (Only keep the variables in the trace set) .

150 151 (4) **Query:** Construct the prompt  $Q = \langle P_{slice}, \Sigma_{proj}^\# \rangle$  and query  $O_{LM}(Q)$ .

152 This process ensures that the LM receives a “minimal sound context”—the smallest subset of  
153 code and data required to mathematically derive the loop invariant.

## 154 6 [CONTRIBUTION 2] OPTIMIZATION VIA INCREMENTAL CACHING

155 While slicing seemingly reduces context pollution, resulting in a shorter, more precise prompt.  
156 There still exist two problems pertaining to runtime efficiency.

157 (1) **Slicing Overhead:** The slicing operation itself can be computationally expensive, especially  
158 for large codebases with complex dependency graphs. Recomputing the slice for every loop  
159 iteration can negate the efficiency gains from context reduction.

160 (2) **Redundant LM Queries:** We did not reduce the number of LM queries. In fact, if it is  
161 possible we made the LM more precise, this may lead to more frequent queries as the  
162 analysis converges slower.

163 I have considered two possible solutions to these problems: *Slice Caching* and Parallelization through Independence.

$$P_{slice} : (\text{FuncID}, \ell, V_{div}) \rightarrow P_{slice} \quad (1)$$

### 168 6.1 Invariant Caching

169 The volatility of the abstract state  $\Sigma^\#$  typically prevents caching LM queries, as the specific values  
170 of variables change in every iteration. However, by asking the LM for *symbolic* invariants (e.g.,  
171 “ $i < 100$ ”) rather than concrete next-steps, we decouple the query from the specific numeric values  
172 in  $\Sigma^\#$ .

173 We define a canonicalization function  $\kappa(Q)$  that normalizes variable names and abstracts concrete  
174 values.

$$\text{Cache}_{LM} : \text{Hash}(\kappa(P_{slice})) \rightarrow \text{Invariant}$$

175 If the structural logic of the loop slice remains unchanged, we retrieve the cached invariant,  
176 bypassing the LM entirely. This effectively reduces the asymptotic complexity of the analysis from  
177  $O(\text{queries} \times \text{latency})$  to  $O(1)$  for previously analyzed code paths.

## 180 7 EVALUATION

### 181 7.1 Experimental Setup

182 We evaluate Slice-AbsInt against two baselines: a traditional Pure Abstract Interpreter (standard  
183 widening) and a Full-Context LM-Guided Abstract Interpreter (a slightly modified lean implemen-  
184 tation of AbsInt-AI). We implement all three analyzers in Lean 4, building on the LeanJavaScript  
185 framework for JavaScript analysis. The evaluation aims to answer the following research questions:

- 186 • **RQ1 (Precision):** Does removing context pollution allow the analyzer to reject infeasible  
187 paths and false positives?
- 188 • **RQ2 (Semantics):** Does a focused context enable the LLM to identify more complex  
189 invariants (e.g., exponential growth)?
- 190 • **RQ3 (Efficiency):** Does slicing reduce token consumption and analysis time compared to  
191 full-context baselines?

### 193 7.2 Notes on Evaluation

194 Some notes on evaluation that have ultimately affected the results:

- 197 • **LLM Choice:** We use Ollama's Llama 3 model for all LM-guided analyses. While not  
198 specialized for code, Llama 3 demonstrates strong general reasoning capabilities [2]. Future  
199 work could explore code-specific models like Gemini, GPT-5. I do believe that larger models  
200 would perform better with larger contexts, with the downsides of cost. This is something  
201 that should be explored in future work.
- 202 • **Benchmark Suite:** We utilize a subset of the SV-COMP [1] benchmark suite, focusing  
203 on JavaScript programs with complex control flow and data structures. Benchmarks were  
204 selected to highlight scenarios where traditional widening fails due to imprecise invariants.
- 205 • **Variable types:** Currently, the implementation only supports integer variables and interval  
206 abstractions. Extending support to heap abstractions and object properties is left for future  
207 work. Object abstraction is where AbsInt-AI found the most success in reducing false  
208 positives, so this is a significant limitation of the current evaluation.
- 209 • **No Functions or Global State:** The current implementation only supports straight-line  
210 code without functions or global state. This simplification was necessary to focus on the  
211 core slicing mechanism. Future work should extend the framework to handle function  
212 calls, recursion, and global variables. This is another area where I believe the slicing would  
213 provide significant benefits, as functions often introduce a lot of irrelevant context, but it's  
214 semantics are crucial for the LM.
- 215 • **Graph Representation vs Syntactic Representation:** Currently, I have only evaluated the  
216 syntactic representation of the slice. Future work should compare the two representations to  
217 see which yields better results with the LLM. My hypothesis is that the graph representation  
218 would perform better, as it is more concise and directly captures dependencies without  
219 syntactic noise.

### 220 7.3 Test 1: Modular Independence

221 To test the functionality of the slicer, and act on my hypothesis that context pollution degrades LM  
222 performance, I create a test program containing two independent modules:

- 223 • **Module A: Noise:** An exponential backoff loop (timeout) with complex arithmetic.
- 224 • **Module B: Target:** A simple linear counter (retries).

225 Our slicer should identify that Module A is irrelevant when targeting the `retries` variable. The  
226 full program and sliced program are shown in Figure ??.

#### 227 A. Original Program (Full Context)

```
228
229 // Initialization
230 timeout := 1;
231 retries := 0;
232
233
234 // MODULE A: Noise
235 // (Exponential Complexity)
236 while timeout < 2000 do
237     timeout := timeout + timeout;
238
239 // MODULE B: Target
240 // (Linear Complexity)
241 while retries < 5 do
242     retries := retries + 1;
```

#### 243 B. Sliced Program (Target: `retries`)

```
// timeout init removed
retries := 0;

// MODULE A: Pruned
// ( identified as dead code )
skip;

// MODULE B: Target
// (Linear Complexity)
while retries < 5 do
    retries := retries + 1;
```

246    7.3.1 *Results.* The results of the modular independence test are summarized in Table 1.

247

248

249    Table 1. Comparison of Analysis Results for Modular Independence (Test W11)

Analyzer	Noise Loop (timeout)	Target Loop (retries)	Outcome
Pure Abstract Interpreter	$[1, \infty]$	$[0, \infty]$	Sound (Baseline)
Full-Context LLM	$[-\infty, +\infty]$	$[-\infty, +\infty]$	<b>Polluted (Panic)</b>
<b>SAINT (Sliced)</b>	$[1, +\infty]$	$[0, \infty]$	<b>Clean (Sound)</b>

255

256

257    7.3.2 *Note to Professor.* Looking at these results, I believe that slicing is working as intended. What  
 258    I am struggling with is building a more precise Abstract Interpreter. It shouldn't be the case that  
 259    we should widen at 5, maybe narrowing needs to be implemented. I've played around with my  
 260    widening thresholds, and anything above 5 loops causes it to hang. If I had more time I would focus  
 261    on improving the underlying abstract interpreter to be more precise, as I believe that would allow  
 262    the LM to make better predictions.

263

## 264    8 CONCLUSION

### 265    8.1 Further Work

266    This work presents an interesting opportunity for further exploration in Abstract Interpretation  
 267    with External libraries. External Libraries and API's, make Abstract Interpretation particularly  
 268    challenging, as the semantics of these functions are often unknown or too complex to model  
 269    precisely. Future work could explore how LLM calls can be integrated to summarize the effects  
 270    of external library calls, potentially using slicing to isolate the relevant parts of the program that  
 271    interact with these libraries.

272    In terms of further work on this paper, there is a lot to be done. First, the abstract interpreter  
 273    needs to be more robust, supporting heap abstractions, functions, and global state. Second, the  
 274    evaluation needs to be more comprehensive, exploring a wider range of benchmarks and LLMs.  
 275    Finally I need to build the caching mechanisms, in order to reduce the analysis runtime.

276

### 277    8.2 What I Learned

278    I found that really building the static analyzer set everything in perspective for me. Watching all  
 279    the theories and the rules, turn into computation was really rewarding. I also learned a lot about  
 280    prompt engineering for LLMs, and how to really think about what context is necessary for the  
 281    model to make good predictions, and how to get as precise of an answer as possible. I am definitely  
 282    excited to play with Lean more in the future. It is such an interesting language.

283

### 284    8.3 Accomplishments

285    In this paper I have experimented on the idea of context pollution in neurosymbolic static analysis,  
 286    and proposed a novel slicing mechanism to mitigate it. I have built a prototype static analyzer in  
 287    Lean 4, implementing the core slicing algorithm and integrating it with an LLM oracle. I have also  
 288    conducted preliminary evaluations demonstrating the effectiveness of slicing in improving analysis  
 289    precision.

290

### 291    ACKNOWLEDGMENTS

292    TBD

293

294

**295 REFERENCES**

- 296 [1] Dirk Beyer. 2011. Software Verification Competition (SV-COMP). In *Tools and Algorithms for the Construction and*  
297 *Analysis of Systems (TACAS)*. Springer, 505–518.
- 298 [2] Mark Chen, Jerry Tworek, Heewoo Jun, et al. 2021. Evaluating Large Language Models Trained on Code. *arXiv preprint*  
299 *arXiv:2107.03374* (2021).
- 300 [3] Patrick Cousot and Radhia Cousot. 1977. Abstract Interpretation: A Unified Lattice Model for Static Analysis of  
301 Programs by Construction or Approximation of Fixpoints. In *Proceedings of the 4th ACM SIGACT-SIGPLAN Symposium*  
302 *on Principles of Programming Languages (POPL)*. ACM, 238–252.
- 303 [4] Alex Wang, Chen Liu, and Yi Zhang. 2025. AbsInt-AI: Sound Static Analysis via Neurosymbolic Guidance. In *Proceedings*  
304 *of the 52nd ACM SIGPLAN Symposium on Principles of Programming Languages (POPL)*. ACM, New York, NY, USA.
- 305 [5] Mark Weiser. 1984. Program Slicing. *IEEE Transactions on Software Engineering SE-10*, 4 (1984), 352–357.
- 306
- 307
- 308
- 309
- 310
- 311
- 312
- 313
- 314
- 315
- 316
- 317
- 318
- 319
- 320
- 321
- 322
- 323
- 324
- 325
- 326
- 327
- 328
- 329
- 330
- 331
- 332
- 333
- 334
- 335
- 336
- 337
- 338
- 339
- 340
- 341
- 342
- 343